

Real World Fuzzing

Charlie Miller

Independent Security Evaluators

October 19, 2007

cmiller@securityevaluators.com

Agenda

- Fuzzing 101
- Common Fuzzing Problems
- Code Coverage
- Examples

Fuzzing

- A form of vulnerability analysis and testing
- Many slightly anomalous test cases are input into the target application
- Application is monitored for any sign of error



Example

- Standard HTTP GET request
 - § GET /index.html HTTP/1.1
- Anomalous requests
 - § AAAAAA...AAAA /index.html HTTP/1.1
 - § GET //////////index.html HTTP/1.1
 - § GET %n%n%n%n%n%n.html HTTP/1.1
 - § GET /AAAAAAAAAAAAAAAAA.html HTTP/1.1
 - § GET /index.html HTTTTTTTTTTTTTTTTP/1.1
 - § GET /index.html HTTP/1.1.1.1.1.1.1.1
 - § etc...

Different Ways To Fuzz

- Mutation Based - “Dumb Fuzzing”
- Generation Based - “Smart Fuzzing”
- Evolutionary


Mutation Based Fuzzing

- Little or no knowledge of the structure of the inputs is assumed
- Anomalies are added to existing valid inputs
- Anomalies may be completely random or follow some heuristics
- Requires little to no set up time
- Dependent on the inputs being modified
- May fail for protocols with checksums, those which depend on challenge response, etc.
- Examples:
 - § Taof, GPF, ProxyFuzz, etc.




Generation Based Fuzzing

- Test cases are generated from some description of the format: RFC, documentation, etc.
- Anomalies are added to each possible spot in the inputs
- Knowledge of protocol should give better results than random fuzzing
- Can take significant time to set up
- Examples
 - § SPIKE, Sulley, Mu-4000, Codenomicon



Evolutionary Fuzzing

- Attempts to generate inputs based on the response of the program
- Autodafe
 - § Prioritizes test cases based on which inputs have reached dangerous API functions
- EFS
 - § Generates test cases based on code coverage metrics (more later)
- This technique is still in the alpha stage :)



The Problems With Fuzzing

- Mutation based fuzzers can generate an infinite number of test cases... When has the fuzzer run long enough?
- Generation based fuzzers generate a finite number of test cases. What happens when they're all run and no bugs are found?
- How do you monitor the target application such that you know when something “bad” has happened?

The Problems With Fuzzing

- What happens when you find too many bugs?
Or every anomalous test case triggers the same (boring) bug?
- How do you figure out which test case caused the fault?
- Given a crash, how do you find the actual vulnerability
- After fuzzing, how do you know what changes to make to improve your fuzzer?
- When do you give up on fuzzing an application?

Example 1: PDF

- Have a PDF file with 248,000 bytes
- There is one byte that, if changed to particular values, causes a crash
 - § This byte is 94% of the way through the file
- Any single random mutation to the file has a probability of .00000392 of finding the crash
- On average, need 127,512 test cases to find it
- At 2 seconds a test case, thats just under 3 days...
- It could take a week or more...

Example 2: 3g2

- Video file format
- Changing a byte in the file to 0xff crashes QuickTime Player 42% of the time
- All these crashes seem to be from the same bug
- There may be other bugs “hidden” by this bug
- FIXED SILENTLY LAST WEEK - NEVER MIND!

Code Coverage

- Some of the answers to these questions lie in *code coverage*
- Code coverage is a metric which can be used to determine how much code has been executed.
- Works for source code or binaries, although almost all the literature assumes you have source

Line Coverage

- Measures how many lines of code (source code lines or assembly instructions) have been executed.

Branch Coverage

- Measures how many branches in code have been taken (conditional jmps)

```
if( x > 2 )  
    x = 2;
```

- The above code can achieve full line coverage in one test case (ex. $x=3$)
- Requires 2 test cases for total branch coverage (ex. $x=1$, $x=2$).

Path Coverage

- Measures the number of paths executed

```
if ( a > 2 )  
    a = 2;  
if ( b > 2 )  
    b = 2;
```

- Requires

- § 1 test case for line coverage

- § 2 test cases for branch coverage

- § 4 test cases for path coverage

- i.e. $(a, b) = \{ (0, 0), (3, 0), (0, 3), (3, 3) \}$

Path Coverage Issues

- In general, a program with n “reachable” branches will require 2^n test cases for branch coverage and 2^n test cases for path coverage

§ Umm....there's a lot of paths in a program!

- If you consider loops, there are an infinite number of paths
- Some paths are *infeasible*

```
if (x > 2)
    x = 2;
if (x < 0)
    x = 0;
```

§ You can't satisfy both of these conditionals, i.e. there is only three paths through this code, not four

Getting Code Coverage Data

- If you've got source
 - § Instrument the code while compiling
 - gcov
 - Insure++
 - Bullseye

Getting Code Coverage Data

- If you live in the real world
 - § Use Debugging info
 - Pai Mei
 - § Virtualization
 - Valgrind
 - Bochs
 - Xen?
 - § Dynamic code instrumentation
 - DynamoRIO
 - Aprobe

Problems with Code Coverage

- Code can be covered without revealing bugs

```
mySafeCpy(char *dst, char* src){  
    if(dst && src)  
        strcpy(dst, src);  
}
```

- Error checking code mostly missed (and we don't particularly care about it)

```
ptr = malloc(sizeof(blah));  
if(!ptr)  
    ran_out_of_memory();
```

- Only “attack surface” reachable
 - § i.e. the code processing user controlled data
 - § No easy way to measure the attack surface

Now the Examples

- Note: we start with some source code examples but move on to binary only

The Hello World of Code Coverage

- Simple program with 3 paths

```
int main(int argc, char *argv[]){  
    if(argc == 2){  
        if(strstr(argv[1], "hi")){  
            printf(" Hello world\n");  
        }  
    } else {  
        printf("Wrong number of arguments\n");  
    }  
    return 1;  
}
```

Gcov

- Compile with “coverage” flags

```
gcc -g -fprofile-arcs -ftest-coverage -o hello hello.c
```

- This generates a .gcno file for each object file which contains static information about it, such as locations of branches, names of functions, etc

Under the Hood

```
0x00001b0a <main+0>:    push    ebp
0x00001b0b <main+1>:    mov     ebp,esp
0x00001b0d <main+3>:    push    ebx
0x00001b0e <main+4>:    sub     esp,0x14
0x00001b11 <main+7>:    call    0x2ffc <__i686.get_pc_thunk.bx>
0x00001b16 <main+12>:   cmp     DWORD PTR [ebp+8],0x2
0x00001b1a <main+16>:   jne     0x1b77 <main+109>
0x00001b1c <main+18>:   lea     eax,[ebx+0x158a]
0x00001b22 <main+24>:   add     DWORD PTR [eax],0x1
0x00001b25 <main+27>:   adc     DWORD PTR [eax+4],0x0
```

- Additional code added to binary
- 64-bit global variable stores coverage information
- Dumped to disk when `gcov_exit()` is called

Run it

```
$ ./hello there  
$ ./hello hi_there  
Hello world
```

- When you run the program, code coverage information is stored in .gcda files for each object file
- To process these files, run gcov

```
$ gcov hello.c  
File 'hello.c'  
Lines executed:83.33% of 6  
hello.c:creating 'hello.c.gcov'
```

hello.c.gcov

```
-:      0:Source:hello.c
-:      0:Graph:hello.gcno
-:      0:Data:hello.gcda
-:      0:Runs:2
-:      0:Programs:1
2:      1:int main(int argc, char *argv[]){
2:      2:          if(argc == 2){
2:      3:              if(strstr(argv[1], "hi")){
1:      4:                  printf(" Hello world\n");
-:      5:              }
-:      6:          } else {
#####:      7:              printf("Wrong number of arguments\n");
-:      8:          }
2:      9:          return 1;
-:     10:}
```

In June 2007...

- A group of cunning, good looking researchers hacked the iPhone
- How'd we find the bug?



- Fuzzing + Code Coverage!

WebKit

- Most Apple Internet applications share the same code, WebKit
- WebKit is an open source library
- Source code is available via svn:
§ svn checkout <http://svn.webkit.org/repository/webkit/trunk> WebKit



Thanks

- From the development site:

The JavaScriptCore Tests

If you are making changes to JavaScriptCore, there is an additional test suite you must run before landing changes. This is the Mozilla JavaScript test suite.

- So we know what they use for unit testing
- Let's use code coverage to see which portions of code might not be as well tested

Icov

- One problem with gcov is the data is stored in many different files
- Icov is an open source software package which collects data from a whole project and displays it in a nice html report
- It can be a minor pain in the ass to get to work...



Build and Run WebKit

- Build it:

```
WebKit/WebKitTools/Scripts/build-webkit -coverage
```

- Run the test suite:

```
WebKitTools/Scripts/run-javascriptcore-tests -coverage
```

- Add a bunch of stupid links for lcov...sigh :(

- Collect coverage data

```
lcov --directory WebKitBuild/JavaScriptCore.build/Release/  
JavaScriptCore.build/Objects-normal/i386 -c -o testsuite.info
```

- Generate HTML report

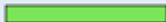


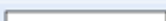


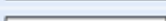

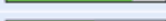
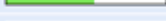

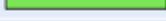






```
genhtml -o WebKit-html -f testsuite.info
```

Results

LTP GCOV extension - code coverage report

Current view: [directory](#)
Test: [testsuite.info](#)
Date: 2007-06-01
Code covered: 59.3 %

Instrumented lines: 13622
Executed lines: 8073

| Directory name | Coverage | | |
|---|---|---------|-------------------|
| /System/Library/Frameworks/CoreFoundation.framework/Headers |  | 100.0 % | 1 / 1 lines |
| /System/Library/Frameworks/JavaVM.framework/Headers |  | 0.0 % | 0 / 53 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/API |  | 0.0 % | 0 / 474 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/bindings |  | 0.0 % | 0 / 530 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/bindings/c |  | 0.0 % | 0 / 190 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/bindings/jni |  | 0.0 % | 0 / 890 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/bindings/objc |  | 0.0 % | 0 / 476 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/kjs |  | 79.3 % | 5723 / 7219 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/pcre |  | 54.7 % | 1338 / 2445 lines |
| /Users/cmiller/woot/WebKit/JavaScriptCore/wtf |  | 0.0 % | 0 / 56 lines |
| /usr/include |  | 100.0 % | 2 / 2 lines |
| /usr/include/architecture/i386 |  | 100.0 % | 3 / 3 lines |
| /usr/include/c++/4.0.0/bits |  | 50.0 % | 4 / 8 lines |
| /usr/share |  | 89.7 % | 96 / 107 lines |
| JavaScriptCore/kjs |  | 84.8 % | 357 / 421 lines |
| kjs |  | 0.0 % | 0 / 39 lines |
| wtf |  | 76.9 % | 528 / 687 lines |
| wtf/unicode/icu |  | 100.0 % | 21 / 21 lines |

Generated by: [LTP GCOV extension version 1.5](#)

Results

- 59.3% of 13622 lines in JavaScriptCore were covered
 - § The main engine (53% of the overall code) had 79.3% of its lines covered
 - § Perl Compatible Regular Expression (PCRE) library (17% of the overall code) had 54.7% of its lines covered
- We decided to investigate PCRE further

...The Rest of the Story

- Wrote a PCRE fuzzer (20 lines of perl)
- Ran it on a standalone PCRE parser (pcredemo from the PCRE library)
- We started getting errors like:

`PCRE compilation failed at offset 6: internal error: code overflow.`

- This was good

A Short Digression on iPhone Hacking: - or - How To Write an Exploit by Fuzzing

- Using our evil regular expression, we could crash mobileSafari (which uses Webkit)
- We didn't have a debugger for the iPhone.
- We couldn't compile code for the iPhone
- We did have crash reports which gave register values
- We did have core dumps (after some iPhone modifications)

All Exploits Need...

- To get control (in this case $pc = r15$)
- To find your shellcode
- Q: How can you do this without a debugger?
- A: The same way you find bugs while watching TV: fuzzing



Fuzz to Exploit

- We generated hundreds of regular expressions containing different number of “evil” strings: “[[*]]”
- Sorted through the crash reports
- Eventually found a good one

A “Good” Crash

Thread 2 crashed with ARM Thread State:

| | | | |
|------------------------------------|----------------|-----------------|-----------------------|
| r0: 0x00065000 | r1: 0x0084f800 | r2: 0x00000017 | r3: 0x15621561 |
| r4: 0x00000018 | r5: 0x0084ee00 | r6: 0x00065000 | r7: 0x005523ac |
| r8: 0x0000afaf | r9: 0x00817a00 | r10: 0x00ff8000 | r11: 0x00000005 |
| ip: 0x15641563 | sp: 0x00552358 | lr: 0x30003d70 | pc: 0x3008cbc4 |
| cpsr: 0x20000010 instr: 0xe583c004 | | | |

| | | | |
|------------------------------|------------|---------------------|-------------------------------|
| __text:3008CBC4 | STR | R12, [R3,#4] | |
| __text:3008CBC8 | BXEQ | LR | |
| __text:3008CBCC | | | |
| __text:3008CBCC loc_3008CBCC | | | ; CODE XREF: __text:3008CBA0j |
| __text:3008CBCC | STR | R3, [R12] | |

- Unlinking of a linked list
- r3 and r12=ip are controllable
- Old school heap overflow (gotta love Apple)
- Gives us a “write anywhere” primitive
- Hows it work? Who the hell knows!
- HD Moore, who is an exploit writing genius, would be sad :(

More Fuzzing For Exploitation

- We decided to overwrite a return address on the stack.
- How do you find it? Fuzz!
 - § True fuzzing folks will call this brute forcing and not fuzzing, but either way its easy...

Exception Type: EXC_BAD_INSTRUCTION

...

Thread 2 crashed with ARM Thread State:

| | | | |
|------------------|-----------------------|-----------------|-----------------------|
| r0: 0x00065038 | r1: 0x00000000 | r2: 0x00000a00 | r3: 0x00000001 |
| r4: 0x00065000 | r5: 0x380135a4 | r6: 0x00000000 | r7: 0x005523e4 |
| r8: 0x00000000 | r9: 0x00815a00 | r10: 0x0084b800 | r11: 0x00000000 |
| ip: 0x380075fc | sp: 0x005523d0 | lr: 0x30003e18 | pc: 0x0055ff3c |
| cpsr: 0x20000010 | instr: 0xffffffff | | |

PNG - with source

- libpng-1.2.16
- Used in Firefox, Safari, and Thunderbird (and others)
- <http://www.libpng.org/pub/png/libpng.html>

Build the Source

- `./configure CFLAGS="-g -fprofile-arcs -ftest-coverage"`
- `make` (errors out)
- `gcc -g -fprofile-arcs -ftest-coverage -I. -L/usr/X11R6/lib/ -I/usr/X11R6/include contrib/gregbook/rpng-x.c .libs/libpng12_la-png.o .libs/libpng12_la-pngset.o .libs/libpng12_la-pngget.o .libs/libpng12_la-pngrutil.o .libs/libpng12_la-pngtrans.o .libs/libpng12_la-pngwutil.o .libs/libpng12_la-pngread.o .libs/libpng12_la-pngrio.o .libs/libpng12_la-pngwio.o .libs/libpng12_la-pngwrite.o .libs/libpng12_la-pngrtran.o .libs/libpng12_la-pngwtran.o .libs/libpng12_la-pngmem.o .libs/libpng12_la-pngerror.o .libs/libpng12_la-pngpread.o .libs/libpng12_la-pnggccrd.o contrib/gregbook/readpng.c -o contrib/gregbook/rpng-x -lX11 -lz -lgcov`
- `result: contrib/gregbook/rpng-x`

Quick Test

```
$ ./contrib/gregbook/rpng-x
$ find . | grep gcda
./libs/libpng12_la-png.gcda
./libs/libpng12_la-pngerror.gcda
./libs/libpng12_la-pnggccrd.gcda
./libs/libpng12_la-pngget.gcda
./libs/libpng12_la-pngmem.gcda
./libs/libpng12_la-pngpread.gcda
...
```

How 'bout a Little Dumb Fuzzing Action?

- Grab a PNG off the Internet
 - § The first one I find is from Wikipedia:
PNG_transparency_demonstration_1.png
- Zero out any code coverage data
 - § `lcov --directory . -z`

Generate Some Files

- Use fuzz.c, the “super” fuzzer
 - § Changes 1-17 bytes in each file
 - § New value is random
 - § Does this 8192 times
- The ultimate in dumb fuzzer technology

```
./fuzz > fuzz.out
```

Use the Files

- Use script.sh
 - § Executes the program 10 at a time
 - § Sleeps 5 seconds
 - § Kills any processes
 - § Repeats
 - § Monitors CrashReporter log for crashes

Get Code Coverage

- We covered 10.7% of the lines



```
cp *.c .libs/  
lcov --directory . -c -o fuzz.info  
genhtml -f -o fuzz_html_files fuzz.info  
...
```

- This compares to
 - § 0.4% for getting the usage statement
 - § 745 of 7399 (10.1%) for opening the good file
 - 43 more lines covered by fuzzing...

What's Up?

- That code coverage kinda sucked...
- Did we choose a bad initial file
- Let's try some other files...
 - § Choose 4 other PNG's from the Internet
 - § Fuzz them the same way
 - § Collect data from each separately

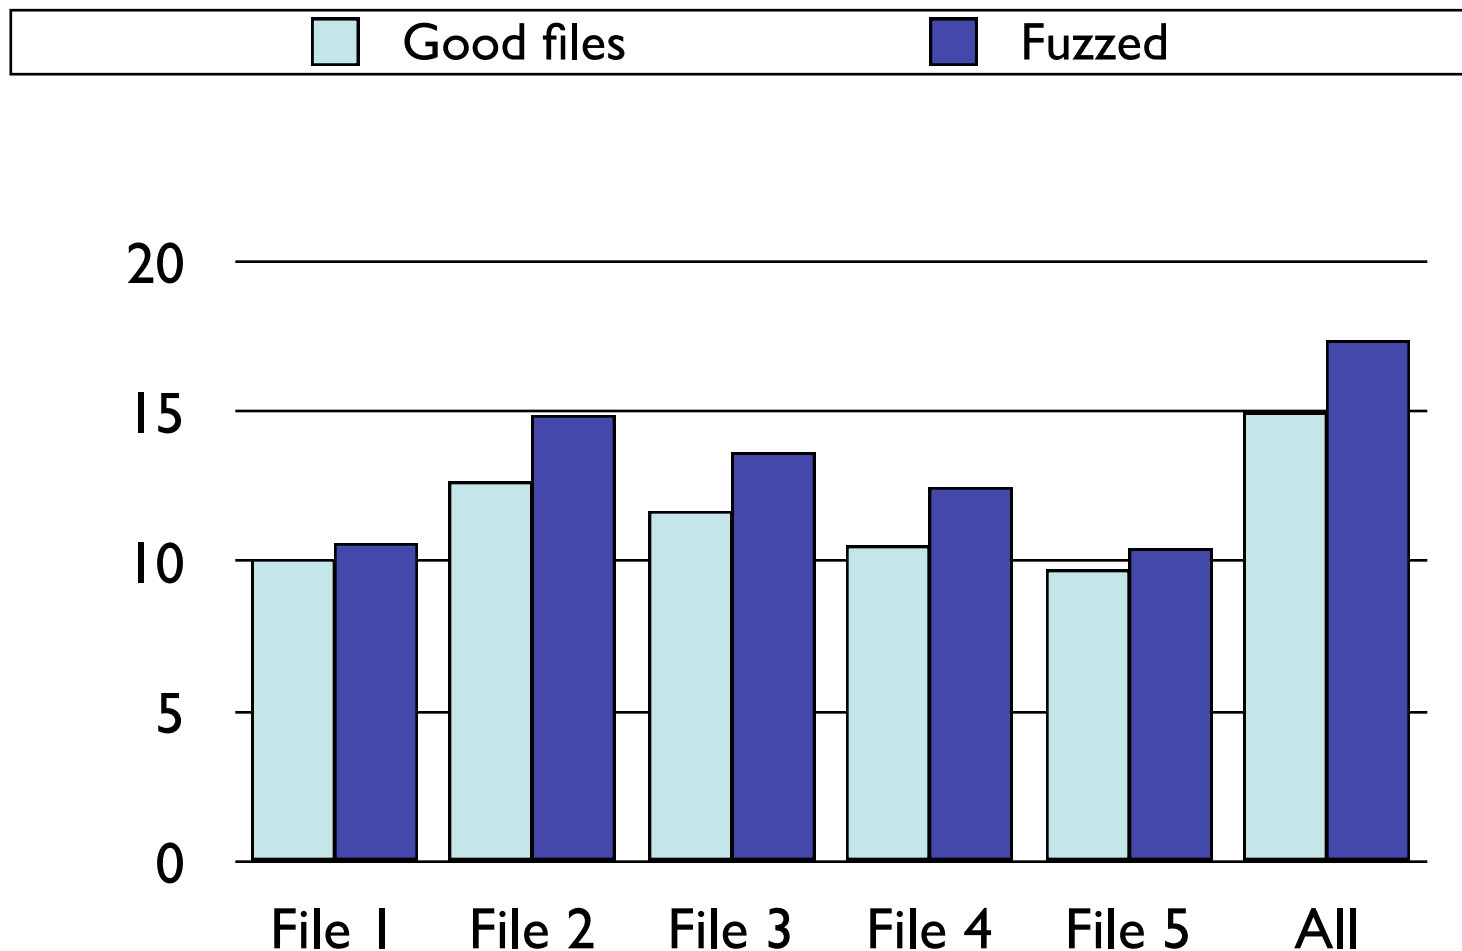
Results



So...

- Initial file can make a big difference
 - § 50% more code coverage from file 2 than in file 5
- What if we ran them all?

The Sum is Greater Than the Parts



WTF is Going On?

- Each PNG contains certain elements that requires some code to process
- Some PNG's contain the same elements, some contain different ones
- By fuzzing with a variety of different PNG's, you increase the chance of having different elements which need processing
- Charlie's Heuristic: Keep adding files until the cumulative effect doesn't increase

A Brief Interlude Into PNG's


- 8 byte signature followed by “chunks”
- Each chunk has
 - § 4 byte length field
 - § 4 byte type field
 - § optional data
 - § 4 byte CRC checksum
- 18 chunk types, 3 of which are mandatory
- Additional types are defined in extensions to the specification
 - § libpng supports 21 chunk types

PNG's From the Wild

- Collected 1631 unique PNG files from the Internet
- Each file was processed and the chunk types present in each was recorded
- Typically, very few chunk types were present

| Number of files | Mean number of chunk types | Standard deviation | Maximum | Minimum |
|-----------------|----------------------------|--------------------|---------|---------|
| 1631 | 4.9 | 1.3 | 9 | 3 |

Distribution of Chunks Found



Observations

- On average, only five of the chunk types are present in a random file!
- 9 of the 21 types occurred in less than 5% of files sampled
- 4 of the chunk types never occurred
- Mutation based fuzzers will typically only test the code from these five chunks
- *They will never fuzz the code in chunks which are not present in the original input*

Enter Generation-Based Fuzzers

- Since Generation-based fuzzers build test cases not from valid data, but from the specification, they should contain all possible chunks
- This should make for a more thorough test


SPIKE

```
//png.spk
// Charlie Miller

// Header - fixed.
s_binary("89504E470D0A1A0A");

// IHDRChunk
s_binary_block_size_word_bigendian("IHDR"); //size of data field
s_block_start("IHDRcrc");
    s_string("IHDR"); // type
    s_block_start("IHDR");
// The following becomes s_int_variable for variable stuff
// 1=BINARYBIGENDIAN, 3=ONEBYE
        s_push_int(0x1a, 1); // Width
        s_push_int(0x14, 1); // Height
        s_push_int(0x8, 3); // Bit Depth - should be 1,2,4,8,16, based
on colortype
        s_push_int(0x3, 3); // ColorType - should be 0,2,3,4,6
        s_binary("00 00"); // Compression || Filter - shall be 00 00
        s_push_int(0x0, 3); // Interlace - should be 0,1
    s_block_end("IHDR");
s_binary_block_crc_word_littleendian("IHDRcrc"); // crc of type and data
s_block_end("IHDRcrc");
...
```

Generation Gap



The Point Is

- Questions such as
 - § How many initial inputs should I use during mutation based fuzzing?
 - § How much effort am I willing to expend to do generational based fuzzing?
- can be answered with code coverage

Halting Problem (Again)


- During all this testing
 - § Used mutation and generation based fuzzers
 - § Generated over 200,000 test cases
 - § ***Not one crash***
- This is a common occurrence for difficult or well audited target applications
- Raises the question: Now what?
- Answer later...
 - § (Hint: has to do with code coverage)



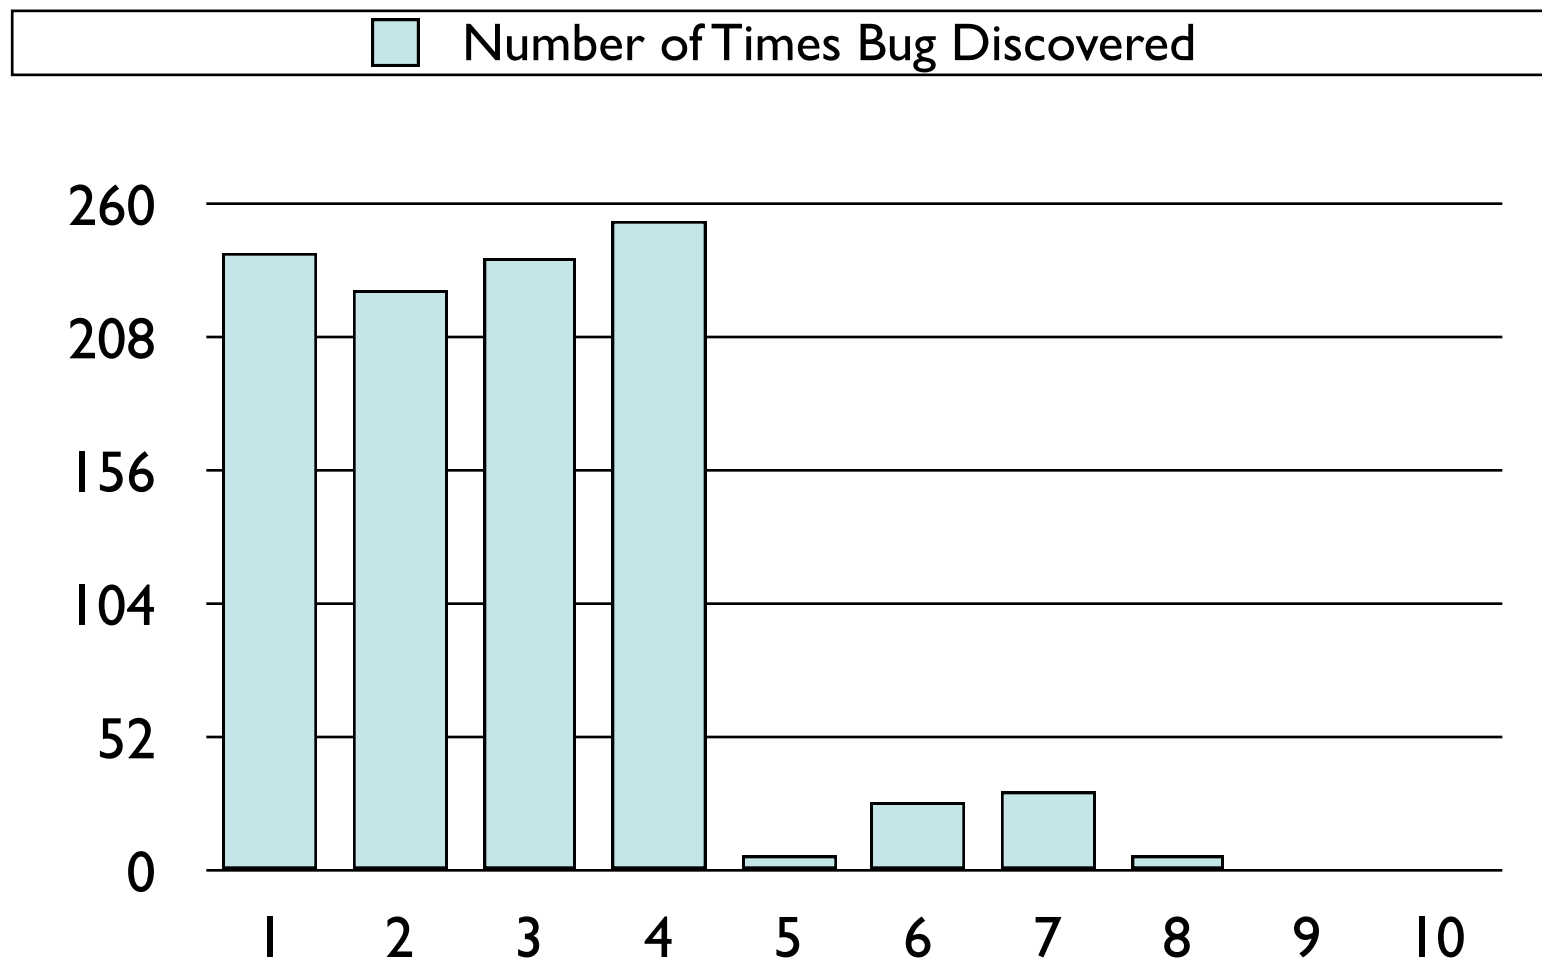
Even More Halting Problem...

- Added 20 “fake” bugs to a server
- Ran ProxyFuzz, a mutation-based fuzzer against it for 450 minutes
- Recorded when each bug was found and how often
- Can you “guess” when to turn off your dumb fuzzer?

Time Required To Find a Bug



Number of Times Discovered



Results

- Sometimes you find a “rare” bug earlier than an “easy” bug
- There are discrete jumps in the time between finding bugs
 - § 4 bugs found in the first 3 minutes
 - Then it took 76 minutes to find the next one
 - § 8 bugs found in the first 121 minutes
 - Then it took another 155 minutes to find the next one
- The final hour didn't find a new bug, what if I would have run it another 24 hours?

Code Coverage is...

- We've seen that code coverage is
 - § A metric to find results about fuzzing
 - § Helpful in figuring out general approaches to fuzzing
 - § Useful to find what code to focus fuzzing upon
- More importantly (we'll see):
 - § A way to improve fuzzing and find more bugs!
 - § Helpful in figuring out when fuzzing is "finished"

Look

- Suppose we didn't know anything about PNG's
- Could we have figured out what was missing when we were fuzzing PNG with the mutation based approach?
- Lets look through some of the Icov report

Yup

```
447      8161 :         else if (!png_memcmp(png_ptr->chunk_name, png_PLTE, 4))
448      0 :             png_handle_PLTE(png_ptr, info_ptr, length);
449      8161 :         else if (!png_memcmp(png_ptr->chunk_name, png_IDAT, 4))
450      :             {
451      8147 :                 if (!png_ptr->mode & PNG_HAVE_IHDR)
452      0 :                     png_error(png_ptr, "Missing IHDR before IDAT");
453      8147 :                 else if (png_ptr->color_type == PNG_COLOR_TYPE_PALETTE &&
454      :                     !png_ptr->mode & PNG_HAVE_PLTE)
455      0 :                     png_error(png_ptr, "Missing PLTE before IDAT");
456      :
457      8147 :                 png_ptr->idat_size = length;
458      8147 :                 png_ptr->mode |= PNG_HAVE_IDAT;
459      8147 :                 break;
460      :             }
461      :         #if defined(PNG_READ_bKGD_SUPPORTED)
462      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_bKGD, 4))
463      0 :                 png_handle_bKGD(png_ptr, info_ptr, length);
464      :             #endif
465      :         #if defined(PNG_READ_cHRM_SUPPORTED)
466      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_cHRM, 4))
467      0 :                 png_handle_cHRM(png_ptr, info_ptr, length);
468      :             #endif
469      :         #if defined(PNG_READ_gAMA_SUPPORTED)
470      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_gAMA, 4))
471      0 :                 png_handle_gAMA(png_ptr, info_ptr, length);
472      :             #endif
473      :         #if defined(PNG_READ_hIST_SUPPORTED)
474      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_hIST, 4))
475      0 :                 png_handle_hIST(png_ptr, info_ptr, length);
476      :             #endif
477      :         #if defined(PNG_READ_oFFs_SUPPORTED)
478      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_oFFs, 4))
479      0 :                 png_handle_oFFs(png_ptr, info_ptr, length);
480      :             #endif
481      :         #if defined(PNG_READ_pCAL_SUPPORTED)
482      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_pCAL, 4))
483      0 :                 png_handle_pCAL(png_ptr, info_ptr, length);
484      :             #endif
485      :         #if defined(PNG_READ_sCAL_SUPPORTED)
486      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_sCAL, 4))
487      0 :                 png_handle_sCAL(png_ptr, info_ptr, length);
488      :             #endif
489      :         #if defined(PNG_READ_pHYs_SUPPORTED)
490      14 :             else if (!png_memcmp(png_ptr->chunk_name, png_pHYs, 4))
491      0 :                 png_handle_pHYs(png_ptr, info_ptr, length);
```

Code Coverage Improves Fuzzing

- Finding spots in the code which are not covered can help with the generation of new test cases
- Beware: covered code doesn't necessarily mean its "fuzzed"
- *Code which has not been executed definitely still needs to be fuzzed!*

Digression into Binary Code Coverage

- So far, we've seen how code coverage can give useful information to help fuzzing
- We've seen how to use gcov and lcov to do this
- The exact same data can be obtained on Windows binaries using Pai Mei
- Pai Mei exists for Mac OS X and is being ported to Linux

Pai Mei

- A reverse engineering framework
- Integrates
 - § PyDbg debugger
 - § IDA Pro databases (via PIDA)
 - § pGraph graphing
 - § mySQL database
- Gives the ability to perform reverse engineering tasks quickly and repeatably
- <http://paimei.openrce.org/>



Pstalker

- A Pai Mei Module
- Uses IDA Pro to get structure of binary
- Sets breakpoints at each basic block (or function)
- Records and removes breakpoints that are hit
- Allows for filtering of breakpoints
- *Gathers code coverage for binaries*

Screenshot

PAIMEIconsole Connections Advanced Help

PAIMEI docs

PAIMEI explore

0xFF
PAIMEI filefuzz

PAIMEI pstalker

Data Sources

Refresh Target List

- Available Targets
 - freedv
 - safari
 - test

PIDA Modules

| # Func | # BB | PIDA Module |
|--------|--------|----------------|
| 7069 | 123325 | coregraphic... |

Add Module(s)

Data Exploration

| # | EIP | TID | Module | Func? | Tag |
|---|----------|------|----------------|-------|------|
| 1 | 6b05b3c0 | 3896 | coregraphic... | Y | test |
| 2 | 6b05b040 | 3896 | coregraphic... | Y | test |
| 3 | 6b05b055 | 3896 | coregraphic... | | test |
| 4 | 6b05b05e | 3896 | coregraphic... | | test |

Functions: 1083 / 7069 Basic Blocks: 11575 / 123325

Data Capture

Refresh Process List

| PID | Process |
|------|-------------------------|
| 280 | toolsrv.exe |
| 1232 | FNPLicensingService.exe |
| 2132 | iPodService.exe |
| 2292 | alg.exe |
| 2668 | wscntfy.exe |
| 3852 | pythonw.exe |
| 3900 | Safari.exe |

Load: Browse

Coverage Depth

☐ Functions

☒ Basic Blocks

☐ Restore BPs ☐ Heavy ☒ Unhandled Or

Start Stalking

```
debugger hit 6b0a33fa cc #11556
debugger hit 6b0a33fc cc #11557
debugger hit 6b0a3417 cc #11558
debugger hit 6b0a341f cc #11559
debugger hit 6b0a3433 cc #11560
debugger hit 6b0a3435 cc #11561
debugger hit 6b0a344d cc #11562
debugger hit 6b0a3470 cc #11563
debugger hit 6b0a3478 cc #11564
debugger hit 6b0a347c cc #11565
debugger hit 6b0d5c85 cc #11566
debugger hit 6b0d5d9e cc #11567
debugger hit 6b0d5dfd cc #11568
debugger hit 6b0d5ce9 cc #11569
debugger hit 6b1dd038 cc #11570
debugger hit 6b1dd62c cc #11571
debugger hit 6b1dd653 cc #11572
debugger hit 6b1dd66c cc #11573
debugger hit 6b1dd6a3 cc #11574
debugger hit 6b1dd6b2 cc #11575
Exporting 11575 hits to MySQL.
Function coverage at 15%. Basic block coverage at 9%.
```

Successfully connected to MySQL server at localhost. Process Stalker

One Hitch

- Can't keep launching the process
- Have to have a way for it to keep loading the fuzzed images
- Just use a meta-refresh tag and point the browser at it

The Fuzzing Website

```
#!/usr/bin/perl


$file = $ENV{'QUERY_STRING'};
$nextfile = $file + 1;
$server = $ENV{'SERVER_NAME'};
$script = $ENV{'SCRIPT_NAME'};
$url = "http://".$server.$script."?". $nextfile;
$pic = sprintf("bad-%d.gif", $nextfile);
$picurl = "http://".$server."/gif/".$pic;

print "Content-type: text/html

<head>
    Fuzz!

";
print " <meta http-equiv=\"refresh\" content=\"2;$url\">";
print " </head><body>";
print "</body>\n";
print "<Script Language=\"JavaScript\">";
print "window.open('$picurl');";
print "</Script>";
```

Missed PNG Basic Blocks




Using Pai Mei to Find the Code

- Do some general browsing in Safari under Pai Mei
 - § Avoid pages with PNG's if possible
 - § Stop when no more breakpoints are hit
- Record this code coverage in a tag
- Filter out on that tag and browse a bunch of different PNG's
- This will record those basic blocks used only in PNG processing (mostly)

This Results In:

- Total basic blocks: 123,325
- Hit during “general browsing”: 12,776
- Hit during PNG only surfing with filter on: 1094 (0.9% of total basic blocks)
 - § This includes 87 functions (out of 7069)
 - § 61 of these basic blocks are in the “main” PNG processing function
 - § Most of the others are in “chunk” specific functions

Where's the Code?



Pai Mei Limitations

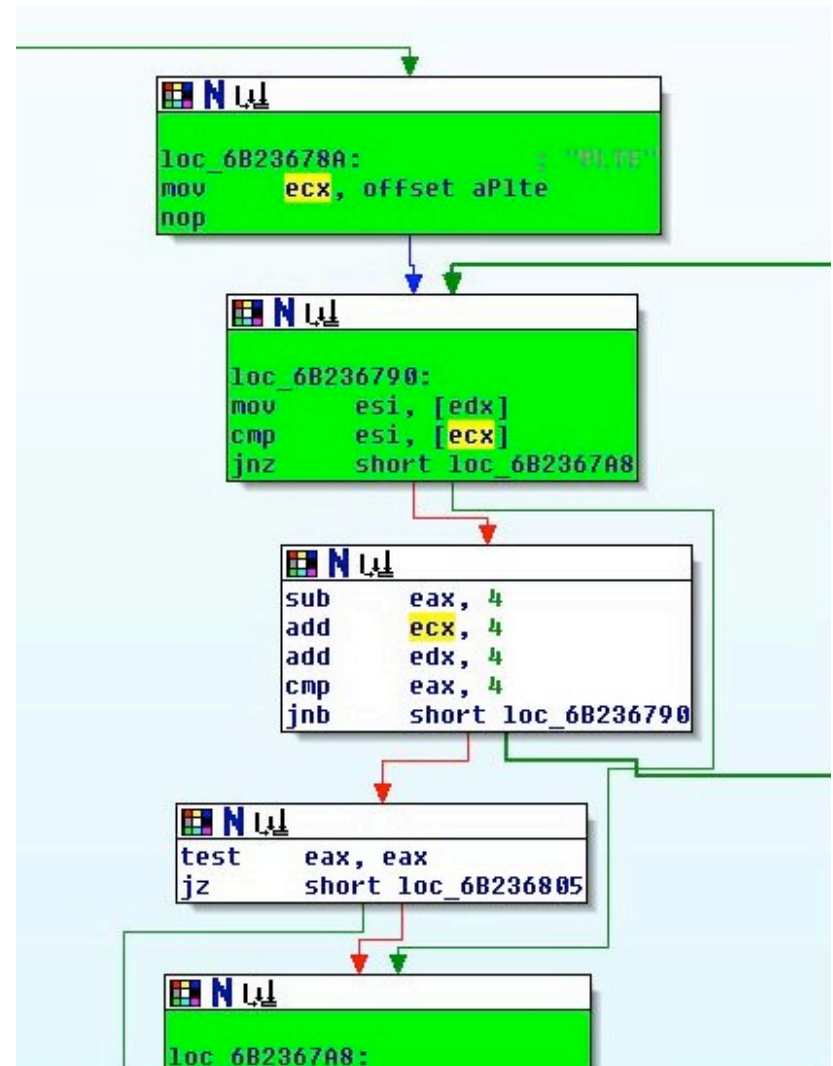
- Pai Mei is only as good as what IDA give it
 - § If IDA misidentifies data as code, bad things happen!
- Some anti-debugging measures screw it up
- Have to know the DLL you're interested in
 - § Or load them all
- For large binaries, it can be slow to set all the breakpoints
 - § For this library, it takes a few minutes

Increasing Code Coverage

- Lack of code coverage is a bad thing
 - § Can't find bugs in code you're not executing
- How do you increase code coverage?
- Basically three ways
 - § Manually
 - § Dynamically using run time information
 - § Automatically from static information

Manually

- You can imagine looking at the PNG code and figuring out how to get more code coverage.



Another Example

- Freeciv 2.0.9, a free multiplayer game similar to Civilization



- Don't ever play this on a computer you care about

Steps to Code Coverage

- Get the Windows binary - no cheating
- Disassemble it
- Dump the PIDA file
- Launch civserver.exe
- Attach with Pai Mei's Pstalker
- Capture a netcat connection to it
- Filter this out (551 of 36,183 BB's - 2%)
- Trace the fuzzing!

GPF

- Great, general purpose mutation-based fuzzer
- Works on packet captures
- Replays packets while injecting faults
- User can manually tell GPF about the structure of the data in the packets
 - § Aids in the anomaly injection
- Many modes of operation

Fuzz FreeCiv

- Start up the game, play a bit
- Capture the packets to a file
- Convert the PCAP file to a GPF file

```
./GPF -C freeciv_reg_game.pcap freeciv_reg_game.gpf
```

- Fire up GPF (main mode)
 - § Main mode replaces some packets with random data


```
./GPF -G 1 ../freeciv_reg_game.gpf client <IP ADDRESS> 5555 ? TCP  
kj3874jff 1000 0 + 0 + 00 01 09 43 close 0 1 auto none
```

FreeCiv Sucks

- Not designed to be fuzzed :)
- Need to add a sleep to GPF so FreeCiv can keep up
- Fuzz overnight...
- I recorded 96 functions during fuzzing
§ 614 / 36183 basic blocks
- Import data back to IDA
- Look for places to increase code coverage

I See One!

- A big switch statement I only hit once



- Tracing back reveals this switch is controlled by the third byte of the packet


Back to GPF

- Up until now we've basically been sending random data
- Using Pai Mei, we observe that the third byte is important
- We modify GPF to make sure it changes the third byte
- We've added a little structure to our random data

```
./bin/GPF -G 1 freeciv_reg_game.gpf client <IP ADDRESS> 5555 ?  
TCP kj3874jfff 1000 0 + 2 2 00 01 255 41 finish 0 1 auto none
```


Better Code Coverage

- 2383 basic blocks covered (after filtering)
 - § Compare this to 614 with the first fuzz run
 - § 4x improvement



- All cases taken in switch (as expected)
- However, still no bugs...

Manual Method Explained

- Send mostly random data
- Examine code coverage to see what structure in the data is important
- Send data which has some elements set but some mostly random parts
- Rinse and Repeat

Fuzzing Beyond the 3rd Byte

- This command replaces the bytes 3 through 10 of each packet, one at a time, with all possible values from 0 to 255

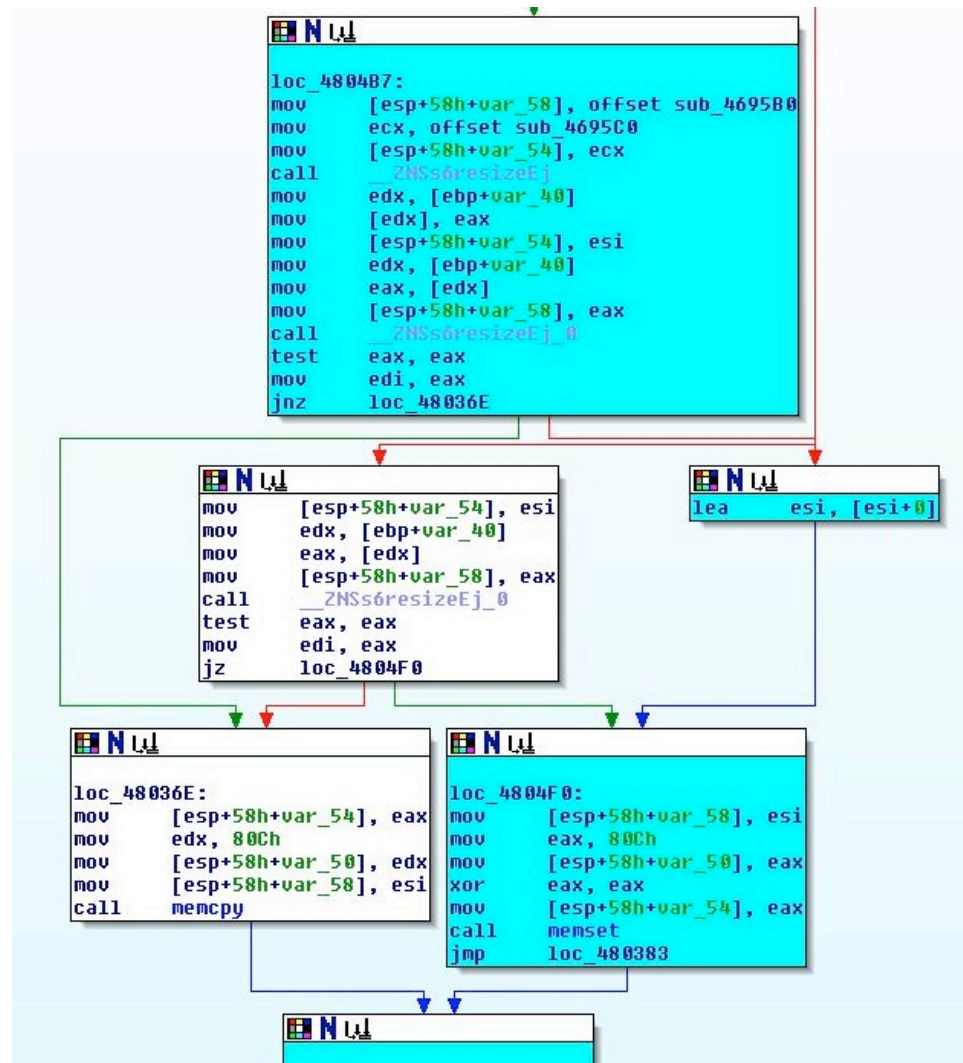
```
./GPF -G 1 ../freeciv_reg_game.gpf client <IP ADDRESS> 5555 ?  
TCP kj3874jff 1000 0 + 2 10 00 01 255 41 finish 0 1 auto none
```

- This will ensure that all the cases in the switch statement are hit and each case will have some random data
- After a bit, CPU is pegged: Memory consumption bug!

Dig Deeper

- Following the methodology, fix the 3rd byte to, say 0x47
- Send in random data to that part of the program
- See what you missed
- Try to do better

Missed Some Spots



Bugs In FreeCiv Aren't a Huge Deal

- Fun for hacking your friends
- Also MetaServer is nice

Freeciv servers around the world

| Host | Port | Version | Patches | State | Players | Topic | Last Update | Players Available |
|---|------|---------|-----------------------------------|------------|---------|----------|-------------|-------------------|
| 88-134-81-104-dynip.superkabel.de | 5555 | 2.0.9 | none | Running | 32 | | 2m | 30 |
| p5B20C598.dip.t-dialin.net | 5560 | 2.0.9 | Warserver - PepServer 0.9.5 devel | Pregame | 0 | NEW GAME | 27s | 0 |
| pool-72-90-153-69.nwrknj.east.verizon.net | 5555 | 2.0.9 | none | Running | 10 | New Game | 2m | 8 |
| pool-72-90-153-69.nwrknj.east.verizon.net | 5556 | 2.0.9 | none | Running | 2 | New Game | 2m | 2 |
| pool-72-90-153-69.nwrknj.east.verizon.net | 5557 | 2.0.9 | none | Running | 9 | New Game | 16s | 9 |
| wsip-70-182-164-206.ks.ks.cox.net | 5555 | 2.0.9 | none | Pregame | 0 | | 39s | 0 |
| wsip-70-184-212-144.om.om.cox.net | 5555 | 2.0.9 | none | Running | 5 | | 2m | 5 |
| ww10.ultico.de | 5551 | 2.0.9 | Warserver - PepServer 0.9.5 devel | Pregame | 0 | NEW GAME | 2s | 0 |
| ww10.ultico.de | 5552 | 2.0.9 | Warserver - PepServer 0.9.5 devel | Game Ended | 5 | NEW GAME | 16s | 5 |
| ww10.ultico.de | 5553 | 2.0.9 | Warserver - PepServer 0.9.5 devel | Pregame | 0 | NEW GAME | 48s | 0 |
| ww10.ultico.de | 5554 | 2.0.9 | Warserver - PepServer 0.9.5 devel | Pregame | 0 | NEW GAME | 49s | 0 |
| ww10.ultico.de | 5555 | 2.0.9 | Warserver - PepServer 0.9.5 devel | Pregame | 0 | NEW GAME | 20s | 0 |

Dynamically Generating Better Test Cases

- Manually improving code coverage is, uh, “time intensive”
- Need to automate the process
- Autodafe kinda does this
- But I prefer another of Jared Demott's tools....

EFS

- Uses Pai Mei Pstalker to record code coverage
- Uses Genetic Algorithms to generate new test cases based on code coverage feedback

Genetic Algorithms

- Technique to find approximate solution to optimization problems
- Inspired by evolutionary biology
 - § Define fitness of an organism (test case)
 - § Must define how to recombine two organisms
 - § Must define how to mutate a single organism
- Lots more complexity but that is the basics

GA example

- $f(x) = -x * (x - 10000)$
- Use “single point crossover” of binary representation of numbers for recombination

```
677  : 00000000000000000000000001010100101
9931 : 00000000000000000000010011011001011
-----
651  : 00000000000000000000000001010001011
```

- Flip a bit 10% of the time for mutation
- Fitness is the value in the function

In Practice


- Running it for a few generations gives

| | | | | | | | |
|------------|------------|------------|------------|-------------|------------|------------|------------|
| 134 | 651 | 485 | 7653 | 1354 | 7654 | 134 | 7302 |
| (1322044) | (6086199) | (4614775) | (17961591) | (11706684) | (17956284) | (1322044) | (19700796) |
| 1354 | 7652 | 134 | 7653 | 7302 | 390 | 1350 | 134 |
| (11706684) | (17966896) | (1322044) | (17961591) | (19700796) | (3747900) | (11677500) | (1322044) |
| 390 | 7302 | 134 | 134 | | 134 | 1350 | 134 |
| (3747900) | (19700796) | (1322044) | (1322044) | 70 (695100) | (1322044) | (11677500) | (1322044) |
| 134 | 134 | 268 | 134 | 1350 | 134 | 134 | 2182 |
| (1322044) | (1322044) | (2608176) | (1322044) | (11677500) | (1322044) | (1322044) | (17058876) |
| 134 | 134 | 134 | 2182 | 1618 | 134 | 134 | 2316 |
| (1322044) | (1322044) | (1322044) | (17058876) | (13562076) | (1322044) | (1322044) | (17796144) |
| 134 | 1618 | 1612 | 132 | 2316 | 134 | 2322 | 1158 |
| (1322044) | (13562076) | (13521456) | (1302576) | (17796144) | (1322044) | (17828316) | (10239036) |

- The optimum value is 5000

GA Approaches the Solution

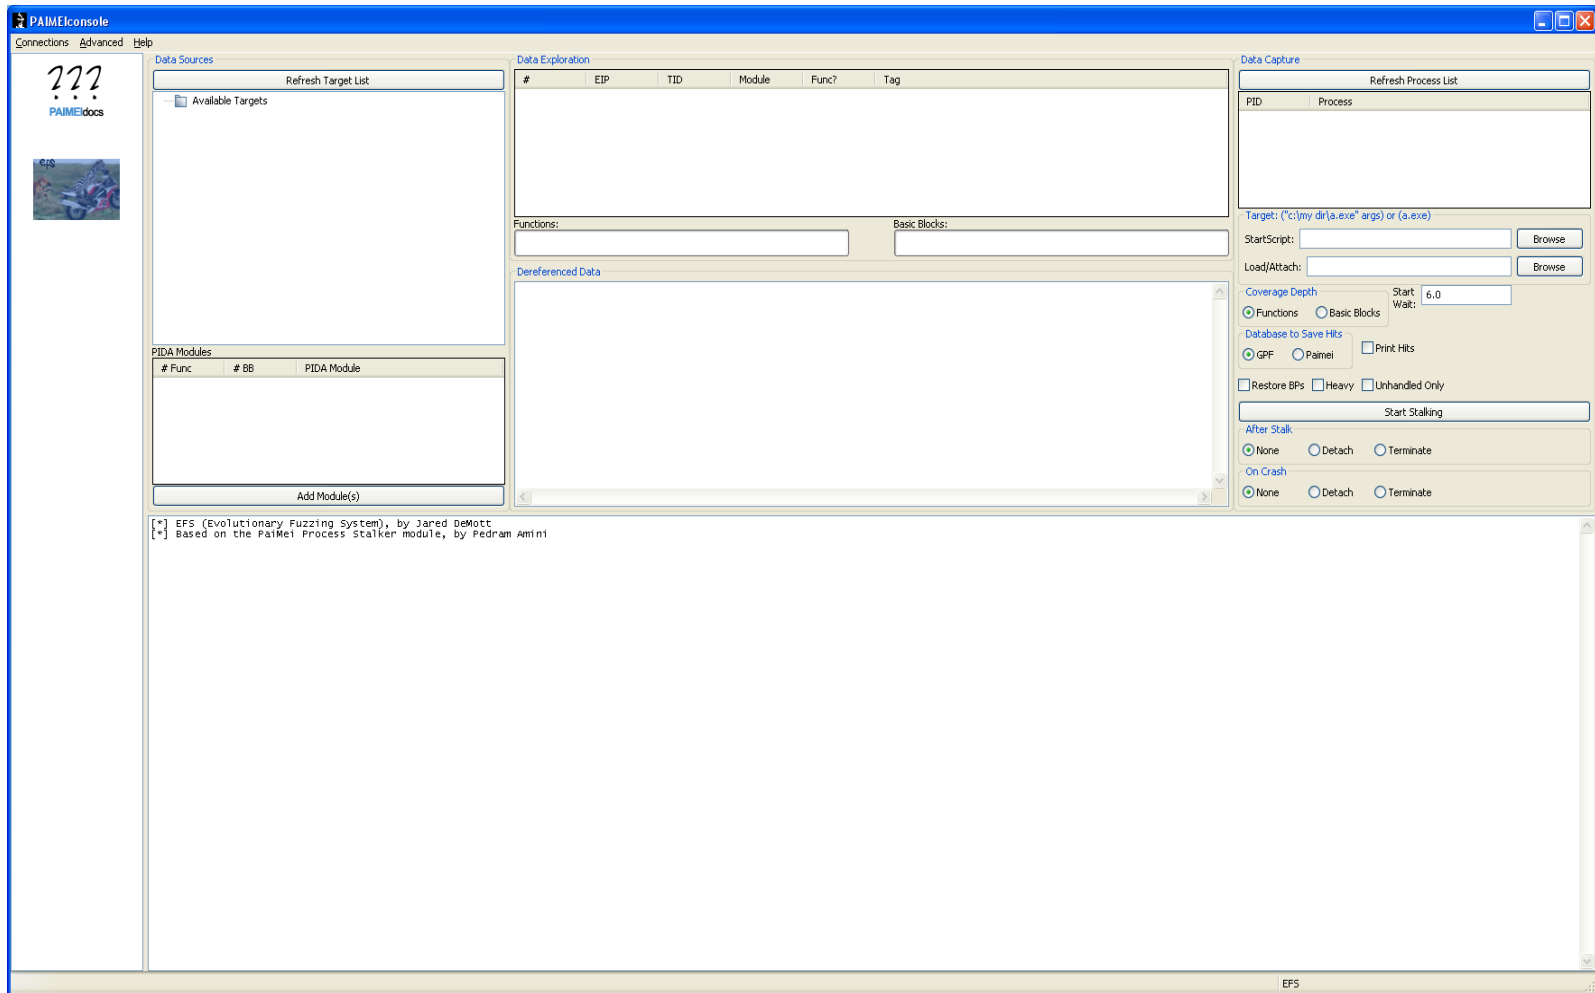
- Generation vs most fit individual
- Approaches the solution



EFS and GA's

- Fitness function: How many functions were covered by the test case (in reality a more elaborate measure is used)
- For breeding, tends to choose the most fit individuals
- Recombination: single point crossover that respects “protocol tokens”
- Mutation: portions of data replaced with fuzzing heuristics

Obligatory Screenshot



Running EFS

- Still needs a PIDA file
- Connect to database
- Add PIDA file to module list
- Enter pathname to application in Load/Attach window
- Choose Connections->Fuzzer Connect
 - § Hit “Listen”
- On Client

```
./GPF -E <IP ADDRESS> root <PASSWORD> 0 0 <IP ADDRESS> 31338 funcs client <IP ADDRESS>  
5555 ? TCP 800000 20 low AUTO 4 25 Fixed 10 Fixed 10 Fixed 35 3 5 9 none none no
```



What You See

```
Successfully played generation 0. Saving to mysqldb.  
Processing Generation 0 ...  
Done processing. Time to play and process: 100 total  
evaluations in  
1001 seconds.  
10.01 sec/eval  
That's 16.683 mins or 0.278 hrs.
```

```
Successfully played generation 1. Saving to mysqldb.  
Processing Generation 1 ...  
Done processing. Time to play and process: 200 total  
evaluations in  
1887 seconds.  
9.44 sec/eval  
That's 31.450 mins or 0.524 hrs.
```

Does It Work?

- The light blue line indicates the most fit pool of testcases
- Code coverage is (slowly) improving



Caveats

- Still experimental
- GA's can get stuck in "local maxima"
- GA's have so many parameters (population size, initial population, mutation percentage, etc), hard to optimize

Statically Generating Code Coverage

- GA's attempt to provide an approximating solution to a difficult problem
- We have the binary, we have the control flow graph, we have the disassembly...
- What if we “solve” the problem exactly?

Existing Work

- Microsoft Research has a tool that generates code coverage maximizing test cases from binaries
 - § <ftp://ftp.research.microsoft.com/pub/tr/TR-2007-58.pdf>
- Catchcov (built on Valgrind) does something similar to try to find integer overflows
- Greg Hoglund has something which tries to do this
- Nothing freely available

General Idea

- Identify where user supplied data enters the program
- Data needs to be traced (symbolically) and branch point's dependence on initial data recorded
- These equations need to be solved, i.e. inputs need to be generated which can go down either branch at each branch point.

Example

- Input comes in through `argv[1]`
- `test()` takes an this value as an int
- 3 possible paths through the program

```
int test(int x){
    if(x < 10){
        if(x > 0){
            return 1;
        }
    }
    return 0;
}

int main(int argc, char *argv[]){
    int x = atoi(argv[1]);
    return test(x);
}
```

Tracing the Data

- Use Valgrind or PyEmu?
- In this trivial example, we'll just do it by hand.
- The constraints would look something like

$$\begin{aligned}x &\geq 10 \\ 0 &< x < 10 \\ x &\leq 0\end{aligned}$$

- In real life, there would be thousands of such constraints

Solve the Constraints

- Can use a Boolean satisfiability solver (SAT)
- One such solver is STP
 - § Constraints expressed as bit vector variables
 - § Bitwise operators like AND, OR, XOR
 - § Arithmetic functions like +, =, *
 - § Predicates like =, <, >

In the STP Language

```
x : BITVECTOR(32);  
QUERY (BVLT (x, 0hex0000000a) );
```

```
x : BITVECTOR(32);  
ASSERT (BVLT (x, 0hex0000000a) );  
QUERY (BVGT (x, 0hex00000000) );
```

```
x : BITVECTOR(32);  
ASSERT (BVLT (x, 0hex0000000a) );  
QUERY (BVLE (x, 0hex00000000) );
```

Solving These Gives

- This gives the test cases $x = \{12, 0, 4\}$
- These give maximal code coverage

```
$ ./stp -p q1
Invalid.
ASSERT( x  = 0hex0000000C  );
$ ./stp -p q2
Invalid.
ASSERT( x  = 0hex00000000  );
$ ./stp -p q3
Invalid.
ASSERT( x  = 0hex00000004  );
```

Using This Technique

- Very sophisticated constraints, such as those that found the Freeciv bug, could be solved (sometimes)
- Optimum test cases can be generated without executing the application
- Combining dynamic and static approaches can optimize fuzzing

Conclusion

- Fuzzing is easy, until you really try it
- Code coverage is a tool that can be used to try to measure and improve fuzzing
- You won't find any bugs in code you haven't tested
- Increasing code coverage can be difficult and time consuming but new tools are coming to make this easier

References

- http://en.wikipedia.org/wiki/Fuzz_testing
- Make My Day: Just Run A Web Scanner, Toshinari Kureha, BH-EU-07
- How Smart is Intelligent Fuzzing - or - How Stupid is Dumb Fuzzing, Charlie Miller, DEFCON 2007
- Robustness Testing Code Coverage Analysis, Teno Rontti, Masters Thesis
- How to Misuse Code Coverage, Brian Marick, <http://www.testing.com/writings/coverage.pdf>
- ProxyFuzz: http://theartoffuzzing.com/joomla/index.php?option=com_content&task=view&id=21&Itemid=40
- STP: <http://theory.stanford.edu/~vganesh/stp.html>
- SPIKE: <http://www.immunitysec.com/downloads/SPIKE2.9.tgz>
- Icov: <http://ltp.sourceforge.net/coverage/lcov.php>
- GPF and EFS: http://www.vdalabs.com/tools/efs_gpf.html

Questions?

- Please contact me at:
cmiller@securityevaluators.com